CNNs in Practical Robotic Applications

Advanced topic I

Dimitris Chrysostomou

Associate Professor

Faculty of Engineering, Department of Materials and Production

Spring 2020

AALBORG UNIVERSITY
DENMARK

Outline

- Previously in OMTP…
- A brief recap of CNNs
- Deep learning and robotic applications
- A glimpse into the deep learning platforms and frameworks
- A couple of practical examples of CNNs
- Assignments
- Recommendations for future endeavors

Previously in OMTP…

Task planning (sequence of skills)

Robot behavior (skills)

Movement primitives (primitives)

Low level control (hardware)

Quick recap of the course

- URDF & XACRO
- GIT and how to maintain your repository
- HRI and DMPs
- Robot control in contact with the environment
- Manipulate things with MoveIT
- Simple object detection and grasping
- Design of behaviors with state machines

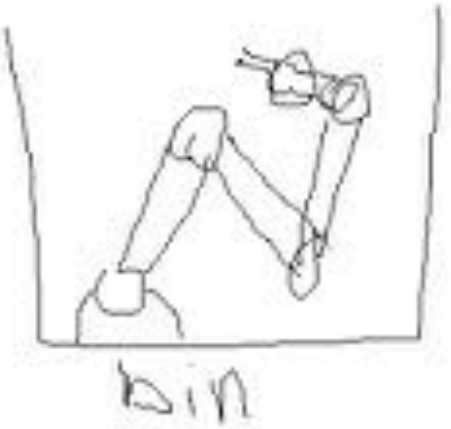

About CNNs and Deep Learning

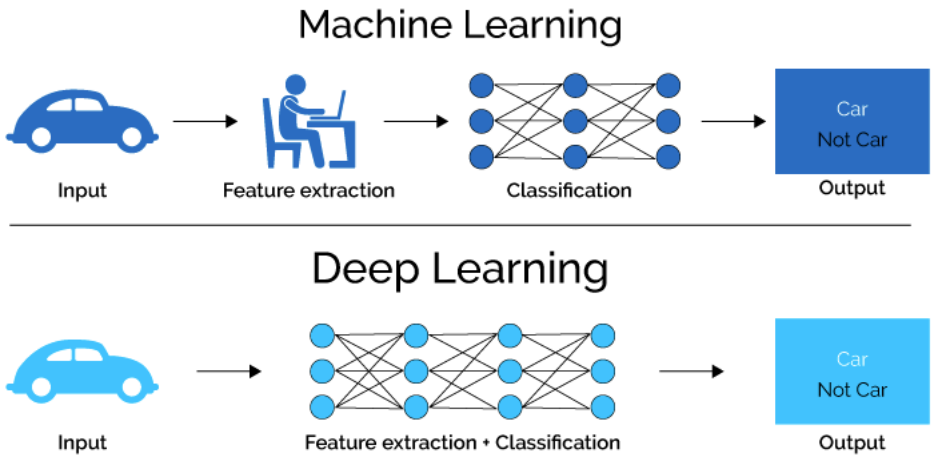

What about Deep Learning?

- You already have some expertise…in different architectures

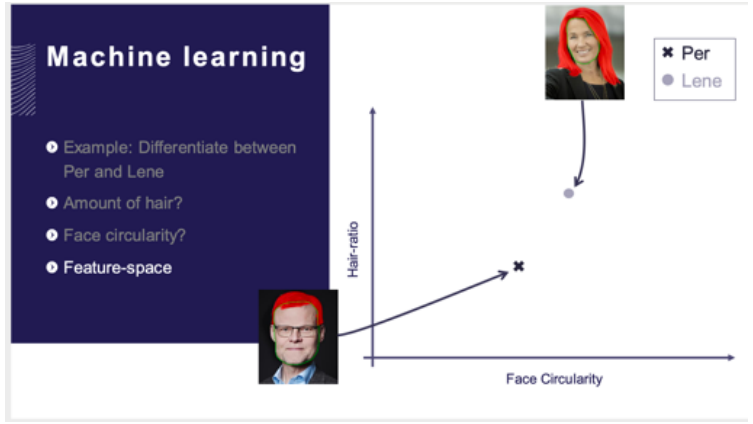

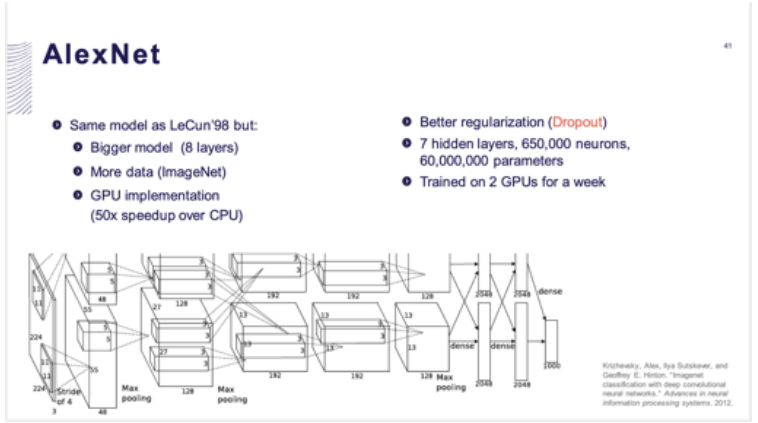

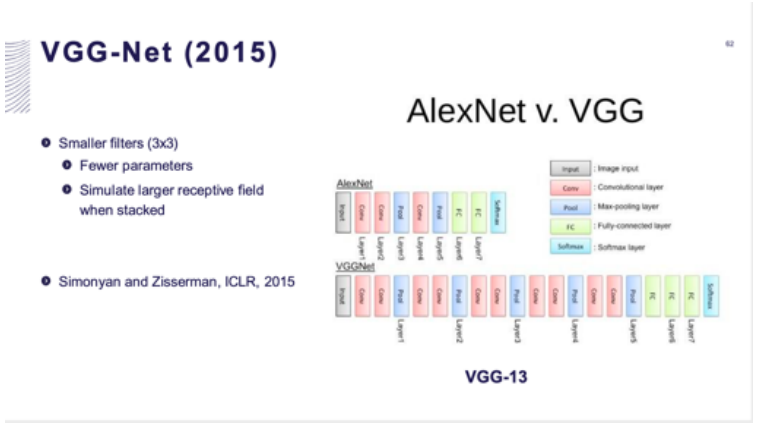

What about Deep Learning?

…CNNs, YOLO, evaluating and understanding training…

Monitoring training

- Train/validation data is typically split 70/30 (or 70/20/10).
- Tune base learning rate and learning rate decay to result in nice learning curve
- Important to monitor train and validation accuracy to avoid overfitting.

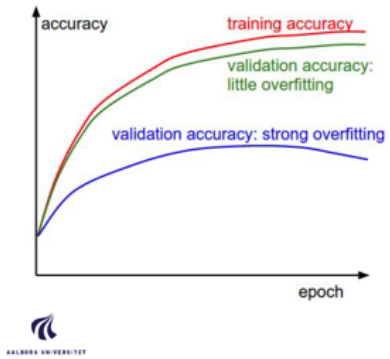

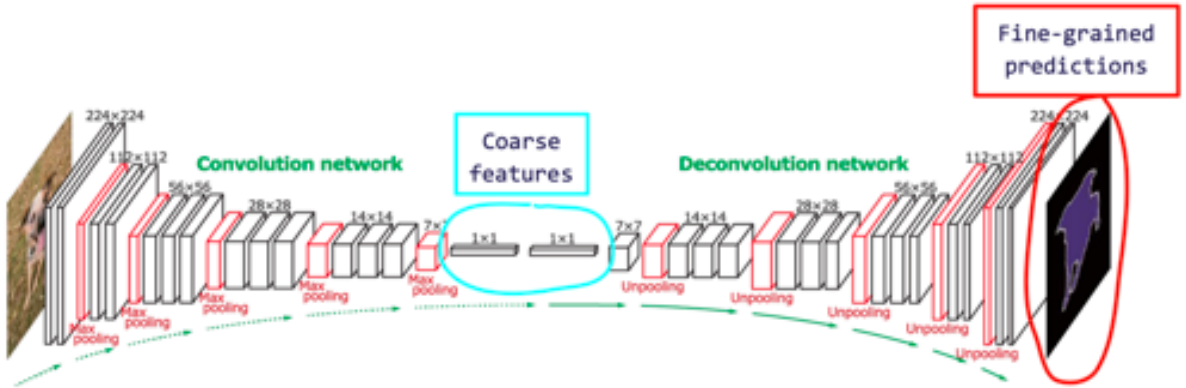

How do we use a neural network?

- Pre-trained network (as we will do today)
- Transfer learning
  - Fine tuning
  - Define new classes, use own training data
- Train a brand new network
  - Learning everything from scratch
  - Requires significant amount of data

What is the fuzz about? GPUs are cheap and plentiful they said…

https://www.youtube.com/watch?v=-P28LKWTzrI

What about Deep Learning and Robotics?

- Deep Learning dominates this field as well
- Robotic applications in every possible field use deep learning
- From predicting the consequences of their possible actions to self-assessing their own competence to complete a task
- Let’s explore some papers

Tossing bot project – Learn to throw stuff

https://www.youtube.com/watch?v=-O-E1nFm6-A

Zeng, A., Song, S., Lee, J., Rodriguez, A. and Funkhouser, T., 2019. Tossingbot: Learning to throw arbitrary objects with residual physics. *arXiv preprint arXiv:1903.11239*. https://arxiv.org/pdf/1903.11239.pdf

Throwing with a robot is VERY difficult!

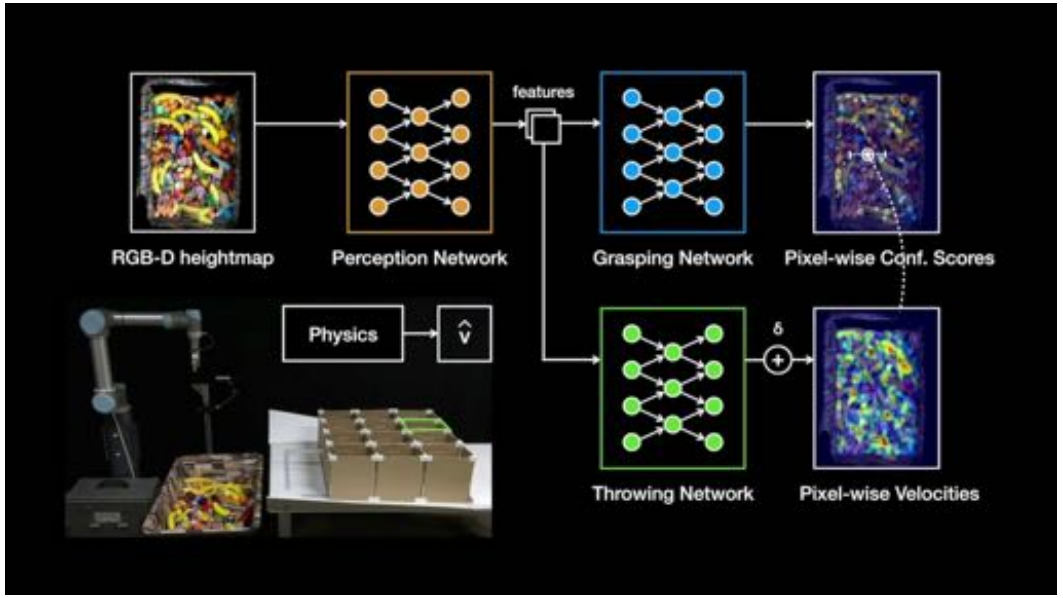

https://www.youtube.com/watch?v=f5Zn2Up2RjQ

End-to-end learning of semantic grasping a.k.a the Google Show-off

https://www.youtube.com/watch?v=iaF43Ze1oeI

Jang, E., Vijayanarasimhan, S., Pastor, P., Ibarz, J. and Levine, S., 2017. End-to-end learning of semantic grasping. *arXiv preprint arXiv:1707.01932*.
https://arxiv.org/pdf/1707.01932.pdf

Learning robotic manipulation

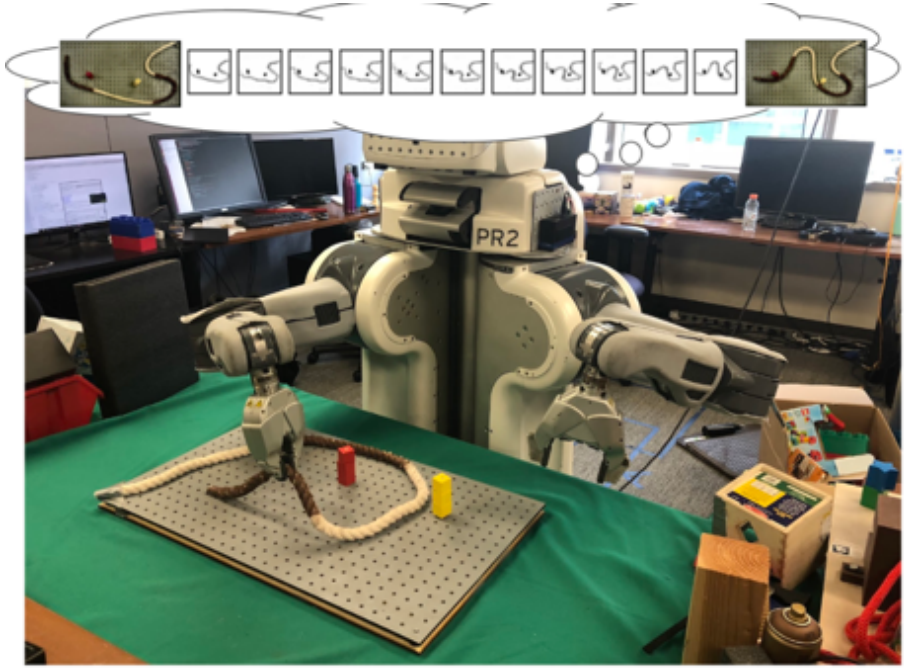

https://www.youtube.com/watch?v=XzL1cf8Esfk

Wang, A., Kurutach, T., Liu, K., Abbeel, P. and Tamar, A., 2019. Learning robotic manipulation through visual planning and acting. arXiv preprint arXiv:1905.04411. https://arxiv.org/pdf/1905.04411 pdf

Translate videos to robotic manipulation

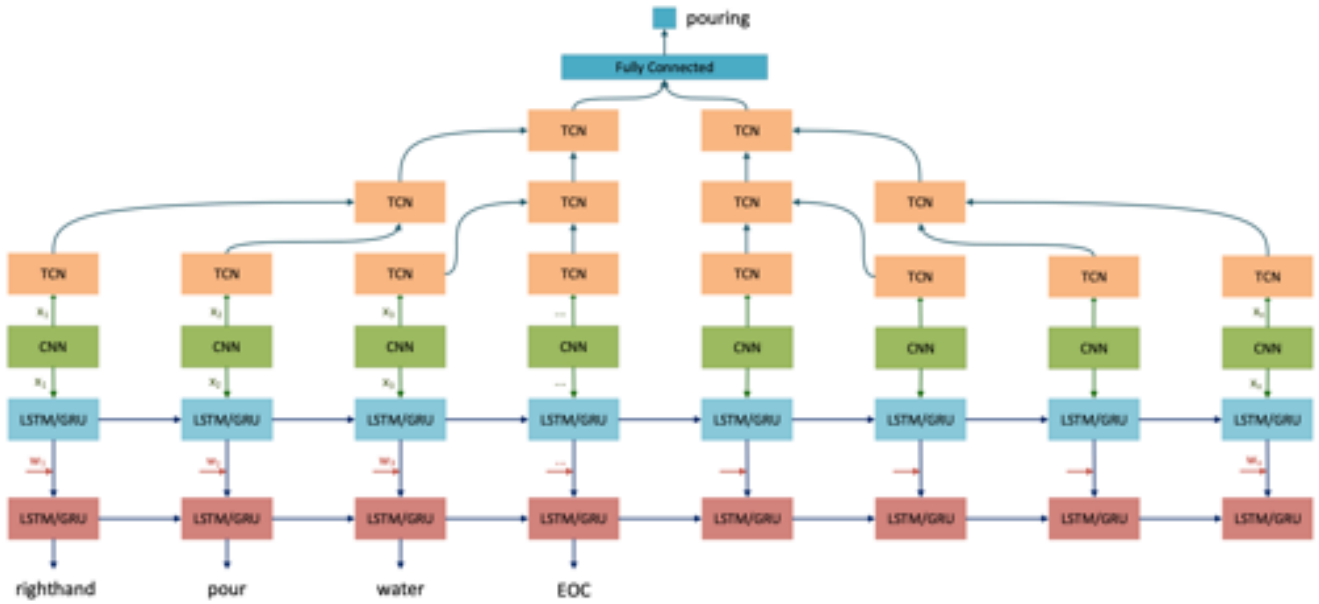

https://www.youtube.com/watch?v=h0Ffcx_Otfo

Nguyen, A., Do, T.T., Reid, I., Caldwell, D.G. and Tsagarakis, N.G., 2019. V2cnet: A deep learning framework to translate videos to commands for robotic manipulation. arXiv preprint arXiv:1903.10869.

Grasp planning and learning grasp policies

DexNet 3.0

https://www.youtube.com/watch?v=tEoTxvD8AeM

DexNet 4.0

https://www.youtube.com/watch?v=r-0PKne9e_w

Mahler, J., Matl, M., Liu, X., Li, A., Gealy, D. and Goldberg, K., 2018, May. Dex-Net 3.0: Computing robust vacuum suction grasp targets in point clouds using a new analytic model and deep learning. In 2018 IEEE International Conference on Robotics and Automation (ICRA) (pp. 1-8). IEEE.

CNN-SLAM: Real-time dense monocular SLAM with learned depth prediction

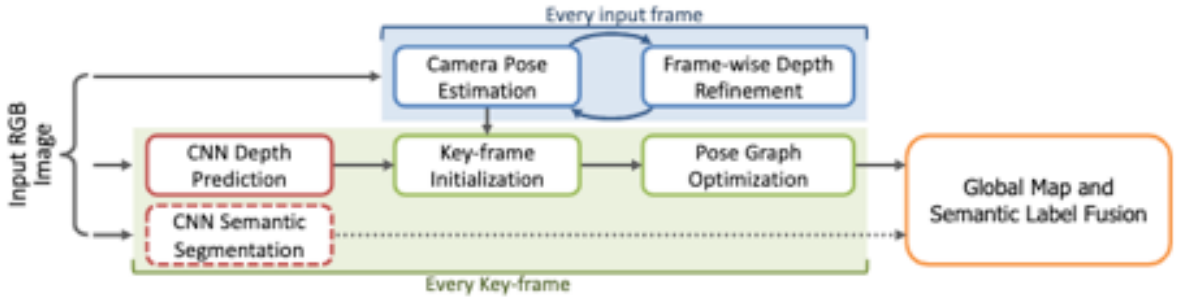

https://www.youtube.com/watch?v=z_NJxbkQnBU

Tateno, K., Tombari, F., Laina, I. and Navab, N., 2017. Cnn-slam: Real-time dense monocular slam with learned depth prediction. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (pp. 6243-6252).

Visual Robot Task planning

Based on learned and artificially generated policies to predict the consequences of future actions

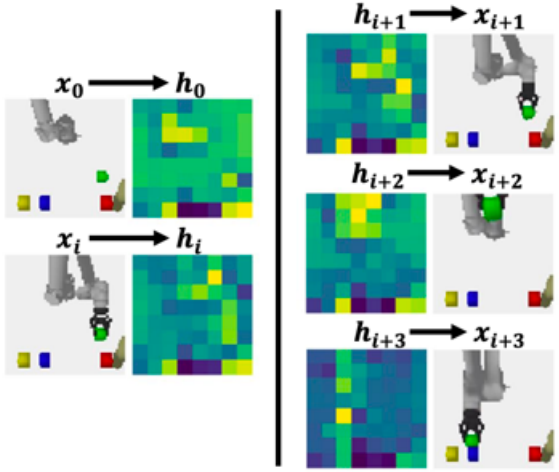

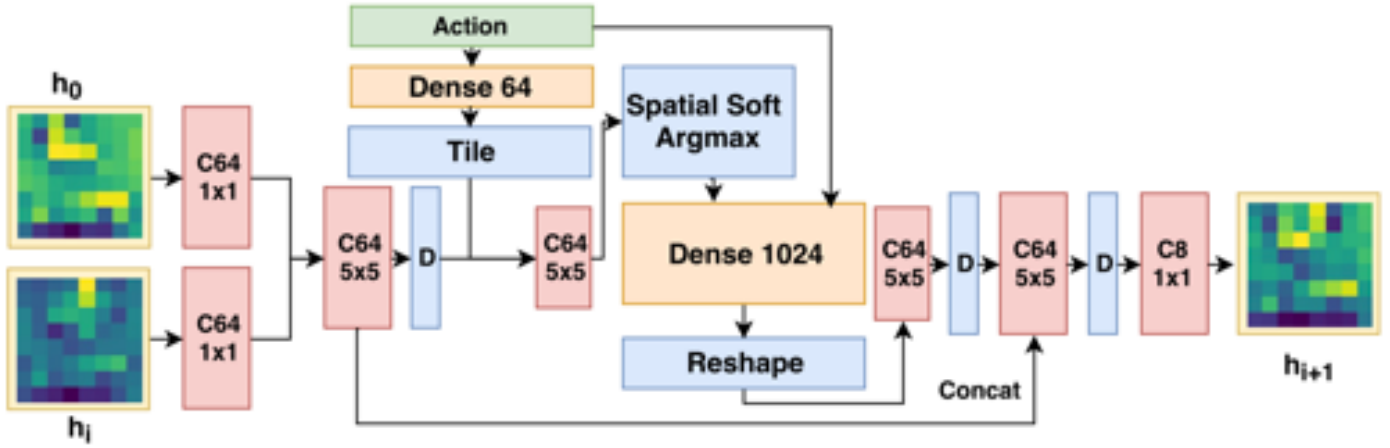

Paxton, C., Barnoy, Y., Katyal, K., Arora, R. and Hager, G.D., 2019, May. Visual robot task planning. In *2019 International Conference on Robotics and Automation (ICRA)* (pp. 8832-8838). IEEE.

Monitoring wildfires with quadcopters

Novac, I., Geipel, K.R., de Domingo Gil, J.E., de Paula, L.G., Hyttel, K. and Chrysostomou, D., 2020, January. A Framework for Wildfire Inspection Using Deep Convolutional Neural Networks. In *2020 IEEE/SICE International Symposium on System Integration (SII)* (pp. 867-872). IEEE.

Instrument segmentation in robot-assisted surgery

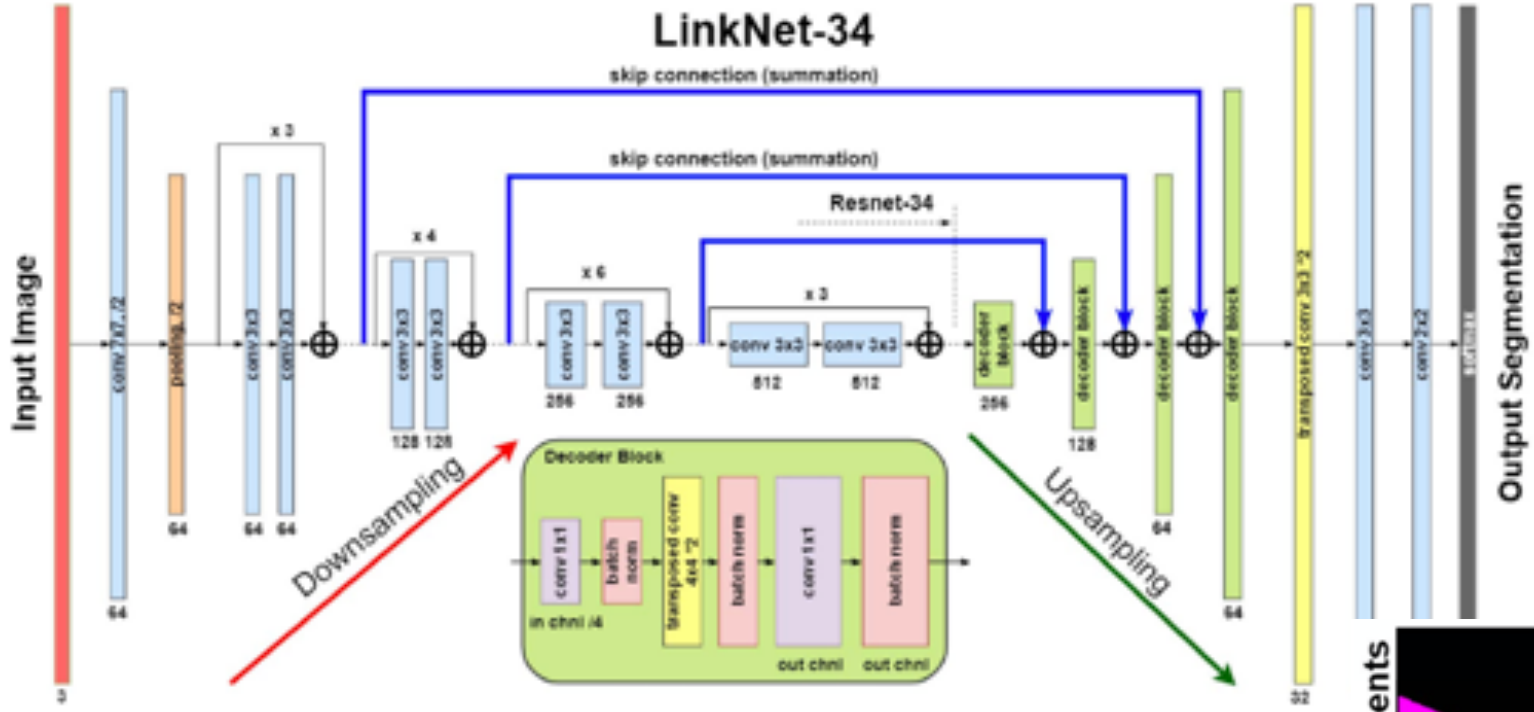

Shvets, A.A., Rakhlin, A., Kalinin, A.A. and Iglovikov, V.I., 2018, December. Automatic instrument segmentation in robot-assisted surgery using deep learning. In *2018 17th IEEE International Conference on Machine Learning and Applications (ICMLA)* (pp. 624-628). IEEE.

Useful literature

- Pierson, H.A. and Gashler, M.S., 2017. Deep learning in robotics: a review of recent research. *Advanced Robotics*, *31*(16), pp.821-835.
- Canziani, A., Paszke, A. and Culurciello, E., 2016. An analysis of deep neural network models for practical applications. *arXiv preprint arXiv:1605.07678*.
- Nassif, A.B., Shahin, I., Attili, I., Azzeh, M. and Shaalan, K., 2019. Speech recognition using deep neural networks: A systematic review. IEEE Access, 7, pp.19143-19165.
- Burghouts, G.J., Huizing, A. and Neerincx, M.A., 2020. Robotic Self-Assessment of Competence. arXiv preprint arXiv:2005.01546.
- Otter, D.W., Medina, J.R. and Kalita, J.K., 2020. A Survey of the Usages of Deep Learning for Natural Language Processing. IEEE Transactions on Neural Networks and Learning Systems.
- Zacharias, J., Barz, M. and Sonntag, D., 2018. A survey on deep learning toolkits and libraries for intelligent user interfaces. arXiv preprint arXiv:1803.04818.
- Esteva, A., Robicquet, A., Ramsundar, B., Kuleshov, V., DePristo, M., Chou, K., Cui, C., Corrado, G., Thrun, S. and Dean, J., 2019. A guide to deep learning in healthcare. Nature medicine, 25(1), pp.24-29.

What about Deep Learning frameworks?

Too many…

Choose your adventure…

Caffe

theano

https://en.wikipedia.org/wiki/Comparison_of_deep-learning_software

…and your online playground

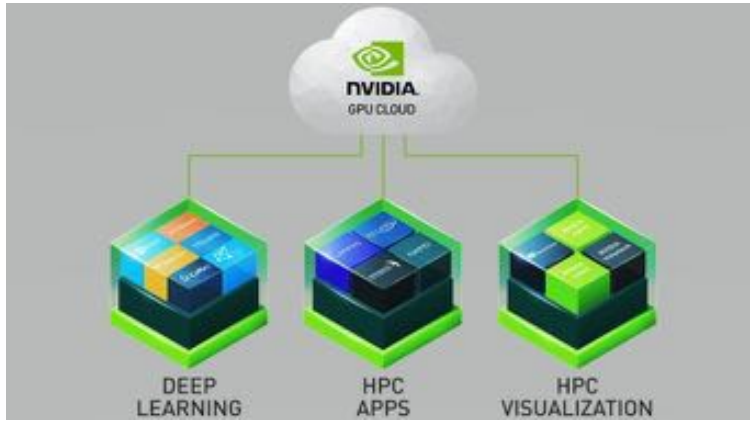

Amazon Machine Learning

colab

Time for some practical applications!

Let’s produce some data

& fight against COVID-19

Let’s produce some data

- What is data augmentation?
  - Answer on Teams

Easy setup

1. From moodle -> Download box_generation.zip -> Unzip
2. Upload the contents of the zip file to your Google drive *(will take some time depending on your network speed)*
3. Login to Google Colab
4. Mount your Google Drive
5. File -> Upload/Open Notebook
6. Select box_generation.ipynb from your Google Drive or from hard drive
7. Edit -> Notebook Settings -> Hardware acceleration -> GPU
8. Run the script and generate new image data

Lots and lots of data

```
1 augment_data("/content/drive/My Drive/augmentation/box_1.png", "/content/drive/My Drive/augmentation/more_data", 5)
```

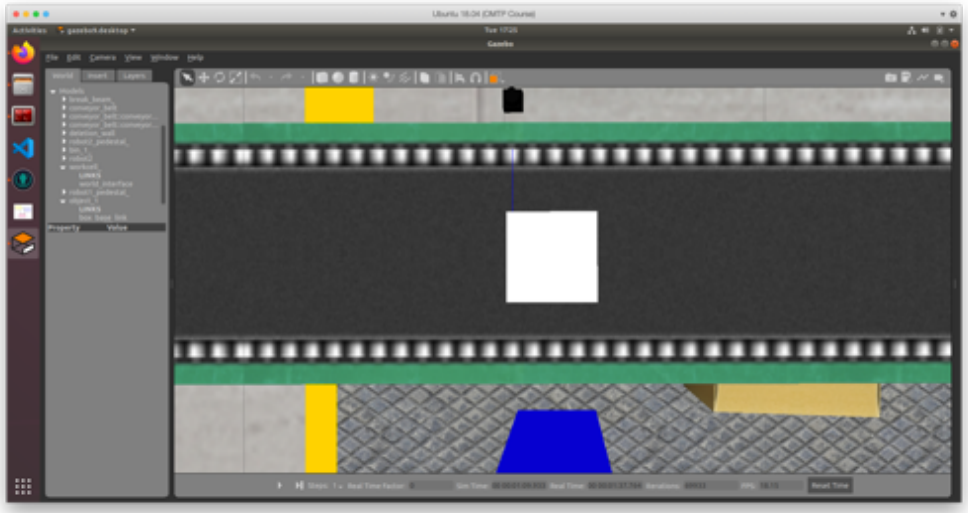

Your original image

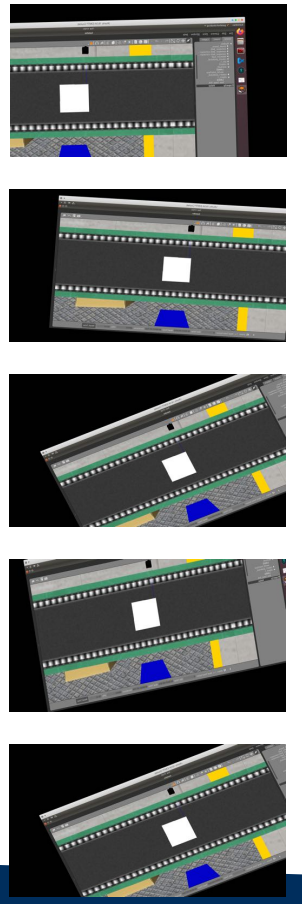

Number of augmented images

Experiment with the parameters

```
# construct the image generator for data augmentation then
# initialize the total number of images generated thus far
aug = ImageDataGenerator(
    rotation_range=30,
    zoom_range=0.15,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.15,
    horizontal_flip=True,
    fill_mode="nearest")

imageGen = aug.flow(image, batch_size=1, save_to_dir=gen_data_path,
    save_prefix="image", save_format="png")
```

What will happen if you play with these parameters?

Hint: https://keras.io/api/preprocessing/image/

TECH | ARTIFICIAL INTELLIGENCE | CORONAVIRUS

France is using AI to check whether people are wearing masks on public transport

The technology won't be used to identify and punish individuals

By James Vincent | May 7, 2020, 10:30am EDT

https://www.theverge.com/2020/5/7/21250357/france-masks-public-transport-mandatory-ai-surveillance-camera-software?_hsenc=p2ANqtz-9L-Z1IQRtr0eXEuDxN6VLANb5tU0fUtvRmSptURnrYwQVvT0V-AoKO2K40cP8yyG2fEclsBkU1zjT0Rvw1QwJa9zB-gg&_hsmi=87877396

World

Paris Tests Face-Mask Recognition Software on Metro Riders

By Helene Fouquet

7 May 2020, 11:46 CEST

https://www.bloomberg.com/news/articles/2020-05-07/paris-tests-face-mask-recognition-software-on-metro-riders

Coronavirus France: Cameras to monitor masks and social distancing

4 May 2020

Coronavirus pandemic

https://tinyurl.com/yare5b2r

Let’s do something useful…

Face Mask Detector with OpenCV, Keras/TensorFlow, & Deep Learning

1. From moodle -> Download covid19_colab.zip -> Unzip
2. Upload the contents of the zip file to your Google drive *(will take some serious time depending on your network speed)*
3. Login to Google Colab
4. Mount your Google Drive
5. File -> Upload/Open Notebook
6. Select covid19_mask.ipynb from your Google Drive or from hard drive
7. Edit -> Notebook Settings -> Hardware acceleration -> GPU

Project structure

- **Sequence**
  - Train_mask_detector.py
  - Detect_mask_image.py
  - Detect_mask_video.py
- **Examples**

It contains three example images for testing the trained model. You can also upload more images to this folder to do your own tests.

```
$ tree --dirsfirst --filelimit 10
.
├── dataset
│   ├── with_mask [690 entries]
│   └── without_mask [686 entries]
├── examples
│   ├── example_01.png
│   ├── example_02.png
│   └── example_03.png
├── face_detector
│   ├── deploy.prototxt
│   └── res10_300x300_ssd_iter_140000.caffemodel
├── detect_mask_image.py
├── detect_mask_video.py
├── mask_detector.model
├── plot.png
└── train_mask_detector.py

5 directories, 10 files
```

Some modifications - Train_mask_detector.py

ANACONDA

PC

colab

```
27 # construct the argument parser and parse the arguments
28 ap = argparse.ArgumentParser()
29 ap.add_argument("-d", "--dataset", required=True,
30   help="path to input dataset")
31 ap.add_argument("-p", "--plot", type=str, default="plot.png",
32   help="path to output loss/accuracy plot")
33 ap.add_argument("-m", "--model", type=str,
34   default="mask_detector.model",
35   help="path to output face mask detector model")
36 args = vars(ap.parse_args())
```

```
27 # The original file uses arguments. Here, they are replaced with the dictionary.
28 # You need to replace the file path where you upload on the google driver.
29 dict = {'dataset': "/content/drive/My Drive/covid19/dataset",
30         'model': "/content/drive/My Drive/covid19/mask_detector.model",
31         'plot': '/content/drive/My Drive/covid19/plot.png'}
```

Some modifications - detect_mask_image.py


```
15 # construct the argument parser and parse the arguments
16 ap = argparse.ArgumentParser()
17 ap.add_argument("-i", "--image", required=True,
18   help="path to input image")
19 ap.add_argument("-f", "--face", type=str,
20   default="face_detector",
21   help="path to face detector model directory")
22 ap.add_argument("-m", "--model", type=str,
23   default="mask_detector.model",
24   help="path to trained face mask detector model")
25 ap.add_argument("-c", "--confidence", type=float, default=0.5,
26   help="minimum probability to filter weak detections")
27 args = vars(ap.parse_args())
```

```
28 # define the path in the dict
29 dict = {'image': "/content/drive/My Drive/covid19/examples/example_02.png",
30         'face': "/content/drive/My Drive/covid19/face_detector",
31         'model': "/content/drive/My Drive/covid19/mask_detector.model",
32         'confidence': 0.5}
```

Some modifications - detect_mask_image.py

Show image

- Since cv2.imshow() is not working in Google Colab, we need to change it
- from google.colab.patches, **“import cv2_imshow”** is added to the “detect_mask_image” script
- Replace cv2.imshow() to cv2_imshow()
- Or just use the edited file from the lecture material…

Some modifications - detect_mask_video.py


```
77 # construct the argument parser and parse the arguments
78 ap = argparse.ArgumentParser()
79 ap.add_argument("-f", "--face", type=str,
80   default="face_detector",
81   help="path to face detector model directory")
82 ap.add_argument("-m", "--model", type=str,
83   default="mask_detector.model",
84   help="path to trained face mask detector model")
85 ap.add_argument("-c", "--confidence", type=float, default=0.5,
86   help="minimum probability to filter weak detections")
87 args = vars(ap.parse_args())
```

```
77 # dictionary for parameters
78 dict = {'face': "/content/drive/My Drive/covid19/face_detector",
79         'model': "/content/drive/My Drive/covid19/mask_detector.model",
80         'confidence': 0.5}
```

Some modifications - detect_mask_video.py

Run the Sequence as is

1. Train_mask_detector.py (~10 mins depending on no. of epochs and local hardware)
2. Detect_mask_image.py
3. Detect_mask_videopy

The original *detect_mask_video.py* cannot be run directly from the Colab.

Since the code is actually running on Google VM, you don't have a camera there but only locally on your computer

There is an extra script to use your webcam to record a realtime video and upload to your Google Drive.

The last part (*detect_mask_video.py*) can detect if you are wearing a mask in the video or not 😷

Some modifications - detect_mask_video.py

```
96 # where the recorded video saved
97 webcam = cv2.VideoCapture("/content/drive/My Drive/covid19/examples/realtimevideo.mp4")
98 fps = int(webcam.get(cv2.CAP_PROP_FPS))
99 # where the prediction results will be saved (images)
100 path = "/content/drive/My Drive/covid19/examples"
```

1. Change the video path - **"/content/drive/My Drive/covid19/examples/realtimevideo.mp4"** to where your video is saved
2. Change the prediction result path – **"/content/drive/My Drive/covid19/examples"** to where you want to save the prediction results.

```
111   # Check if this is the frame closest to 2 seconds
112   if count%(2*fps)==0:
```

3. takes sample image **every 2 seconds** from the video, you can change to other time interval here.

Recap

- Why Deep Learning and CNNs are important
- How CNNs are used in robotics
- What type of Deep Learning frameworks are out there
- Put everything together for an application for a face mask detection in the COVID-19 era

Lecture material

- Play with the code and try different aspects of it. How does it perform?

- Download the papers and keep them for further study

Assignment for all

- Make the code run on your own computer using Google Colab or locally.
- Record a (very short) video of yourself detecting your face…do you wear a mask or not?
- As always, submit assignment to git repo – one short video per team member

Optional assignment #1

- Exchange MobileNetv2 with your favorite network (surprise me!) and show how it performs 🤓
- Bonus question: can we just use whichever network for this application?
- As always, submit assignment to git repo and post on MS Teams

Optional assignment #2

Use it in your OMTP environment and train your simulated camera to identify two classes of (any kind of) objects 😎

Recommendations for future endeavors

- Explore transfer learning
- Always evaluate your performance (IoU, mAP, …)
- Explore TensorFlow and other DL possibilities in ROS
- Choose your territory – Computer Vision, NLP, Speech Recognition, RL – and be an expert in it
- Follow your experts in your field of choice
- Follow the flagship conferences in the field and arXiv.org

- Don’t be too religious on specific frameworks, platforms and/or tools.

https://youtu.be/cJMwBwFj5nQ